

Expected ISO

This project was originally designed for use at the university level and aimed to utilize Trackman data to quantify an expected ISO on contact. Throughout this project, the necessary metrics are labeled with names such as xBA instead of xBAcon, but we are ultimately focusing just on contact. The purpose of this is to create an easy way to pull a quantified “raw power” estimate for any opponent. This makes it easier to know which opposing hitters should be noted for their power potential on our scouting reports.

In order to develop expected ISO, we obviously need to develop expected Batting Average and expected Slugging on contact. In the future, these will be referred to as xBA and xSLG. The first step is to take the complete Trackman dataset from the 2018-2019 season that is available. Unfortunately, this dataset is limited to teams that have entered our data sharing agreement, meaning this sample size is relatively small for this type of analysis. I then subsetting this dataset to only batted ball events and assigned batting average and slugging percentage values for every event.

```
cbb_contact <- subset(cbb_complete, !is.na(cbb_complete$ExitSpeed))
cbb_inplay <- subset(cbb_contact, cbb_contact$PitchCall == 'InPlay')
cbb_inplay$EventID <- c(1:nrow(cbb_inplay))

cbb_inplay$ba <- 0
cbb_inplay$ba[cbb_inplay$PlayResult == 'Single'] <- 1
cbb_inplay$ba[cbb_inplay$PlayResult == 'Double'] <- 1
cbb_inplay$ba[cbb_inplay$PlayResult == 'Triple'] <- 1
cbb_inplay$ba[cbb_inplay$PlayResult == 'HomeRun'] <- 1

cbb_inplay$slg <- 0
cbb_inplay$slg[cbb_inplay$PlayResult == 'Single'] <- 1
cbb_inplay$slg[cbb_inplay$PlayResult == 'Double'] <- 2
cbb_inplay$slg[cbb_inplay$PlayResult == 'Triple'] <- 3
cbb_inplay$slg[cbb_inplay$PlayResult == 'HomeRun'] <- 4
```

I plan to first run a kNN regression based on exit velocity and launch angle, then correct for other factors after. Due to the relatively small size of this data, I’m going to cross-validate the dataset using 10 folds. This next set of code is creating a separate dataframe for each fold, then stitching the original dataframe back together.

```
test_size <- round(nrow(cbb_inplay) * 0.1 ,0)

for(i in 1:10){
  sample_vec <- sample(1:nrow(cbb_inplay), test_size, replace = FALSE)
  sample_frame <- cbb_inplay[sample_vec,]
  cbb_inplay <- cbb_inplay[-sample_vec,]
  assign(paste("test_fold_",i,sep = ''),sample_frame)
}

cbb_inplay <- rbind(test_fold_1, test_fold_2, test_fold_3, test_fold_4, test_fold_5,
```

```

test_fold_6, test_fold_7, test_fold_8, test_fold_9, test_fold_10)

cbb_inplay <- subset(cbb_contact, cbb_contact$PitchCall == 'InPlay')
cbb_inplay$EventID <- c(1:nrow(cbb_inplay))

cbb_inplay$ba <- 0
cbb_inplay$ba[cbb_inplay$PlayResult == 'Single'] <- 1
cbb_inplay$ba[cbb_inplay$PlayResult == 'Double'] <- 1
cbb_inplay$ba[cbb_inplay$PlayResult == 'Triple'] <- 1
cbb_inplay$ba[cbb_inplay$PlayResult == 'HomeRun'] <- 1

cbb_inplay$slg <- 0
cbb_inplay$slg[cbb_inplay$PlayResult == 'Single'] <- 1
cbb_inplay$slg[cbb_inplay$PlayResult == 'Double'] <- 2
cbb_inplay$slg[cbb_inplay$PlayResult == 'Triple'] <- 3
cbb_inplay$slg[cbb_inplay$PlayResult == 'HomeRun'] <- 4

```

Although I know I want to conduct a kNN regression, I don't know what size K is most appropriate. To do so, I am going to run a for loop that runs the analysis at a variety of different values for K. I've nested another for loop inside that will cross-validate the dataset and then keep track of the test errors at each value for K.


```

error_frame <- data.frame("K" = seq(10,100,5),
                          "xBA" = xBA_error,
                          "xSLG" = xSLG_error)

testerror_plot <- ggplot() +
  geom_line(data = error_frame, aes(x = K, y = xBA), color = "blue") +
  geom_line(data = error_frame, aes(x = K, y = xSLG), color = "red") +
  labs(x = 'K', y = 'Test Error', title = 'Test Error by K Size')
testerror_plot

```

Test Error by K Size



The graph indicates a general trend of larger values for K greatly increasing the test error for xSLG, while xBA only starts to noticeably increase towards the end. Since it will be a different algorithm that predicts xBA and xSLG, we could choose two different values for K. However, since the ultimate goal is to predict ISO, I will need to use the same value of K for both predictions in order to prevent the possibility of a negative ISO. Since xBA and xSLG have different scales, I decided to rank their errors at each value for K and find the lowest average of those at a particular value, instead of just adding the errors together.

```
error_frame <- error_frame[order(error_frame$xBA),]
error_frame$ba_rank <- c(1:19)
error_frame <- error_frame[order(error_frame$xSLG),]
error_frame$slg_rank <- c(1:19)
error_frame$rank_avg <- (error_frame$ba_rank + error_frame$slg_rank) / 2
error_frame <- error_frame[order(error_frame$rank_avg, error_frame$ba_rank),]
best_k <- error_frame$K[1]

xba_knn <- knnreg(xba_frame, inplay_train$ba, best_k)
xslg_knn <- knnreg(xslg_frame, inplay_train$slg, best_k)
```

I have now fitted the knn regressions for both xBA and xSLG. In order to check the validity, heatmaps of both statistics are shown below. The shape somewhat resembles similar plots that exist for xWOBA on baseball savant, indicating to me that we are on the right track.

```
ev_vec <- c(round(cbb_inplay$ExitSpeed,0))
la_vec <- c(round(cbb_inplay$Angle,0))
```

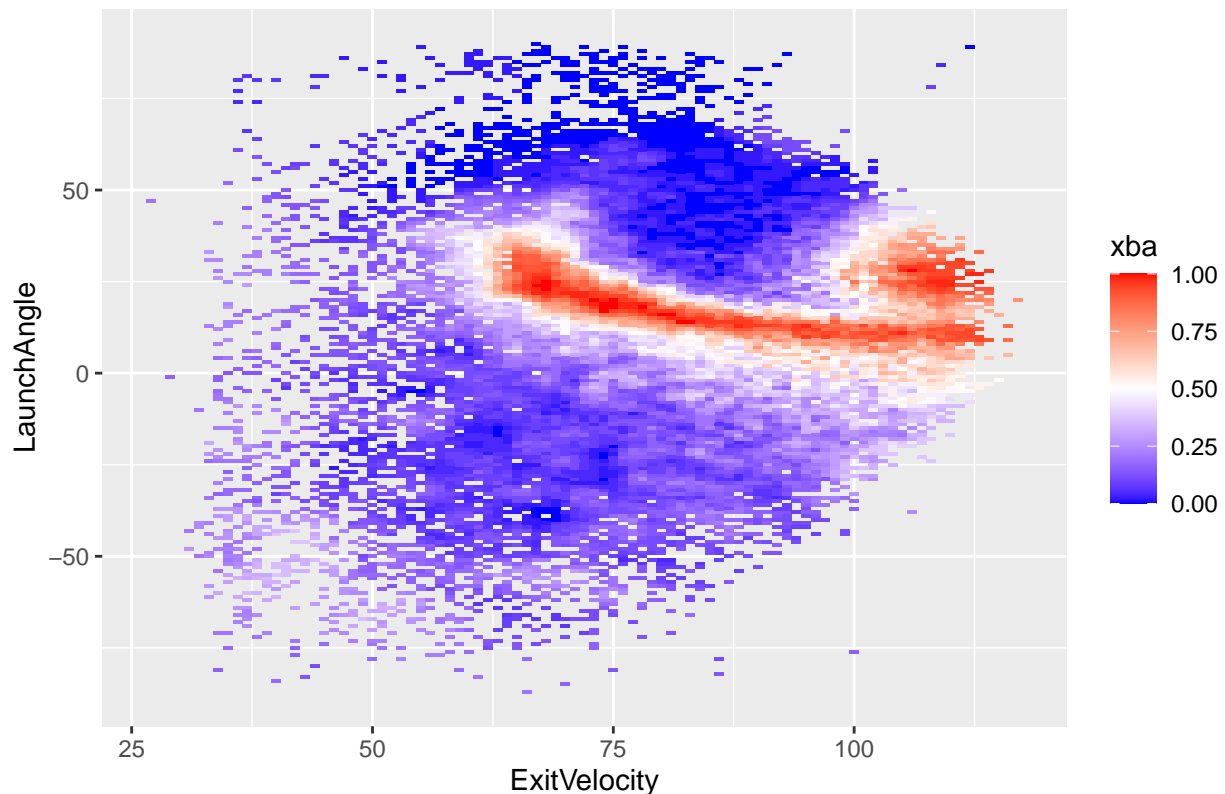
```

expected_display <- data.frame('ExitVelocity' = ev_vec,
                              'LaunchAngle' = la_vec)
expected_display <- unique(expected_display)
display_xba <- predict(xba_knn, expected_display)
display_xslg <- predict(xslg_knn, expected_display)
expected_display$xba <- c(display_xba)
expected_display$xslg <- c(display_xslg)

xba_map <- ggplot(data = expected_display, aes(x = ExitVelocity, y = LaunchAngle)) +
  geom_tile(aes(fill = xba)) +
  scale_fill_gradient2(low = "Blue", mid = "White", high = "Red", midpoint = 0.5, limits = c(0, 1), oob = "white")
ggtitle('College xBA')
print(xba_map)

```

College xBA




```

xslg_map <- ggplot(data = expected_display, aes(x = ExitVelocity, y = LaunchAngle)) +
  geom_tile(aes(fill = xslg)) +
  scale_fill_gradient2(low = "Blue", mid = "White", high = "Red", midpoint = 1, limits = c(0, 2), oob = "white")
ggtitle('College xSLG')
print(xslg_map)

```

College xSLG



Before making further adjustments, I'm going to go ahead and see who the top performers are in the model's current iteration. Remember all data is from 2018 and 2019 and is limited to games that took place at stadiums within our data sharing agreement. The presence of names such as Torkelson, Vaughn, and Bleday near the top indicates that we are on the right track

```
player_frame <- data.frame("ExitSpeed" = cbb_inplay$ExitSpeed,
                           "Angle" = cbb_inplay$Angle)

player_xBA <- predict(xba_knn, player_frame)
player_xSLG <- predict(xslg_knn, player_frame)
player_frame$xBA <- player_xBA
player_frame$xSLG <- player_xSLG
player_frame$BA <- cbb_inplay$ba
player_frame$SLG <- cbb_inplay$slg
player_frame$xISO <- player_frame$xSLG - player_frame$xBA
player_frame$PA <- 1
player_frame$Player <- cbb_inplay$Batter
cbb_inplay$Bearing <- as.character(cbb_inplay$Bearing)
player_frame$Direction <- round(as.numeric(cbb_inplay$Bearing),0)

## Warning: NAs introduced by coercion
```

```
player_averages <- aggregate(cbind(xBA,xSLG,PA)~Player, player_frame, sum)
player_averages$xBA <- player_averages$xBA / player_averages$PA
player_averages$xSLG <- player_averages$xSLG / player_averages$PA
```

```

player_averages$xISO <- player_averages$xSLG - player_averages$xBA

player_averages <- subset(player_averages, player_averages$PA >= 30)
player_averages <- player_averages[order(-player_averages$xISO),]

print(head(player_averages, 20))

```

```

##           Player      xBA      xSLG  PA      xISO
## 433    Cabell, Elijah 0.5217909 1.2035369 56 0.6817460
## 2038   McCann, Kyle 0.4908213 1.1175523 69 0.6267311
## 3229   Vaughn, Andrew 0.5036036 1.0963964 74 0.5927928
## 269    Bleday, JJ 0.4810582 1.0694180 105 0.5883598
## 158    Barr, Cole 0.4978897 1.0713366 57 0.5734469
## 3157 Torkelson, Spencer 0.4912186 1.0602151 124 0.5689964
## 2556   Radcliff, Baron 0.5186380 1.0652330 62 0.5465950
## 254    Bishop, Hunter 0.4864924 1.0261438 102 0.5396514
## 2121   Mendoza, Drew 0.5287689 1.0653729 87 0.5366039
## 1846   Lloyd, Matthew 0.4794118 1.0111111 68 0.5316993
## 3292   Washer, Jake 0.5648746 1.0939068 31 0.5290323
## 528    Chavers, Parker 0.4160643 0.9386881 83 0.5226238
## 2055   McDaniel, Joel 0.4397661 0.9469786 57 0.5072125
## 231    Berryhill, Luke 0.4063492 0.9064713 91 0.5001221
## 2921    Smith, Armani 0.4575866 0.9560335 93 0.4984468
## 2719   Sabato, Aaron 0.4788594 0.9762045 113 0.4973451
## 1471   Horanski, Luke 0.4312618 0.9276836 59 0.4964218
## 1197   Gorski, Matt 0.4961857 0.9810945 67 0.4849088
## 316    Bradley, Scotty 0.4404938 0.9204938 45 0.4800000
## 295    Boone, Trevor 0.4620072 0.9362007 62 0.4741935

```

I will now graph the difference between expected BA and SLG and actual BA and SLG. This is done to see if there are any biases in the model that may be caused by external forces

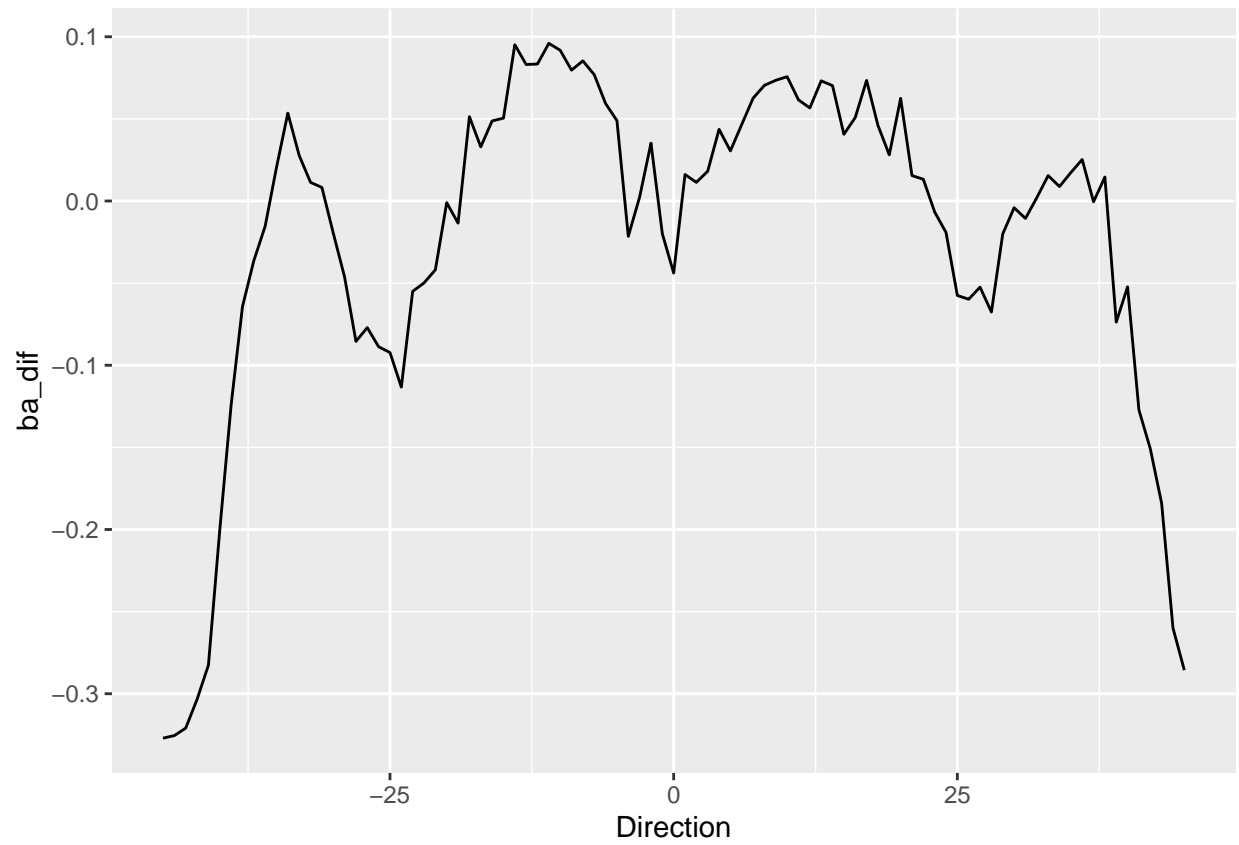
```

direction_averages <- aggregate(cbind(xBA,xSLG,PA,BA,SLG)~Direction, player_frame, sum)
direction_averages$xBA <- direction_averages$xBA / direction_averages$PA
direction_averages$xSLG <- direction_averages$xSLG / direction_averages$PA
direction_averages$xISO <- direction_averages$xSLG - direction_averages$xBA
direction_averages$BA <- direction_averages$BA / direction_averages$PA
direction_averages$SLG <- direction_averages$SLG / direction_averages$PA
direction_averages$ba_dif <- direction_averages$xBA - direction_averages$BA
direction_averages$slg_dif <- direction_averages$xSLG - direction_averages$SLG


direction_averages <- subset(direction_averages, direction_averages$Direction >= -45 &
                             direction_averages$Direction <= 45)

direction_line <- ggplot(data = direction_averages, aes(x = Direction, y = ba_dif)) +
  geom_line()
direction_line

```



```
direction_line2 <- ggplot(data = direction_averages, aes(x = Direction, y = slg_dif)) +  
  geom_line()  
direction_line2
```



As the graphs above indicate, there is a clear bias in the model based on the direction of the batted ball. Given that the positions where xBA outperforms actual BA are at the approximate normal positions for infielders, it would appear defensive positioning is at least partially the cause of this bias.

To correct this, I am going to fit a generalized additive model for both xBA and xSLG. I initially chose the number of splines based on the shape of each graph. Once I ran the GAM's the first time, the new graphs still appeared to show some form of bias, so I added a second term to correct for these. I then applied the adjustments to the xBA and xSLG metrics that already existed from the kNN regression. However, I opted to add no adjustment on batted balls that were outside of the foul lines since these were prone to unrealistic adjustments and the only batted balls included in this data set were marked "In-Play", meaning balls out of the foul lines were limited to easy pop-ups.

```
ba_adj <- gam(ba_dif~ns(Direction,7) + ns(Direction,9), data = direction_averages)
```

```
## Warning in model.matrix.default(mt, mf, contrasts): non-list contrasts argument
## ignored
```

```
summary(ba_adj)
```

```
##
## Call: gam(formula = ba_dif ~ ns(Direction, 7) + ns(Direction, 9), data = direction_averages)
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
```



```
## -0.072463 -0.011368 0.001107 0.013192 0.065540
##
## (Dispersion Parameter for gaussian family taken to be 5e-04)
##
## Null Deviance: 0.95 on 90 degrees of freedom
## Residual Deviance: 0.0359 on 75 degrees of freedom
## AIC: -421.0822
##
## Number of Local Scoring Iterations: 2
##
## Anova for Parametric Effects
##           Df Sum Sq Mean Sq F value    Pr(>F)
## ns(Direction, 7)  7 0.69768 0.099669 208.410 < 2.2e-16 ***
## ns(Direction, 9)  8 0.21643 0.027054  56.571 < 2.2e-16 ***
## Residuals        75 0.03587 0.000478
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
slg_adj <- gam(slg_dif~ns(Direction,6) + ns(Direction,9), data = direction_averages)
```

```
## Warning in model.matrix.default(mt, mf, contrasts): non-list contrasts argument
## ignored
```

```
summary(slg_adj)
```

```
##
## Call: gam(formula = slg_dif ~ ns(Direction, 6) + ns(Direction, 9),
## data = direction_averages)
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -0.167022 -0.024954  0.005965  0.023844  0.168962
##
## (Dispersion Parameter for gaussian family taken to be 0.0026)
##
## Null Deviance: 7.4687 on 90 degrees of freedom
## Residual Deviance: 0.2001 on 78 degrees of freedom
## AIC: -270.6762
##
## Number of Local Scoring Iterations: 2
##
## Anova for Parametric Effects
##           Df Sum Sq Mean Sq F value    Pr(>F)
## ns(Direction, 6)  6 6.9233 1.15388 449.894 < 2.2e-16 ***
## ns(Direction, 9)  6 0.3454 0.05757  22.446 3.469e-15 ***
## Residuals        78 0.2001 0.00256
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
predict_frame <- data.frame("Direction" = as.numeric(as.character(cbb_inplay$Bearing)))
```

```
## Warning in data.frame(Direction = as.numeric(as.character(cbb_inplay$Bearing))):
## NAs introduced by coercion
```

```

player_frame$ba2 <- predict(ba_adj, predict_frame)

## Warning in predict.lm(object, newdata, se.fit, scale = 1, type = if (type == :
## prediction from a rank-deficient fit may be misleading

player_frame$ba2 <- ifelse(player_frame$Direction >= -45 & player_frame$Direction <= 45,
                           player_frame$ba2, 0)

predict_frame <- data.frame("Direction" = as.numeric(as.character(cbb_inplay$Bearing)))

## Warning in data.frame(Direction = as.numeric(as.character(cbb_inplay$Bearing))):
## NAs introduced by coercion

player_frame$slg2 <- predict(slg_adj, predict_frame)

## Warning in predict.lm(object, newdata, se.fit, scale = 1, type = if (type == :
## prediction from a rank-deficient fit may be misleading

player_frame$slg2 <- ifelse(player_frame$Direction >= -45 & player_frame$Direction <= 45,
                           player_frame$slg2, 0)

```

After adding these values, it's time to graph the difference between expected and actual batting average and slugging percentage to see if the GAM appeared to correct the bias.

```


player_frame$xBA <- player_frame$xBA - player_frame$ba2
player_frame$xSLG <- player_frame$xSLG - player_frame$slg2

direction_averages <- aggregate(cbind(xBA,xSLG,PA,BA,SLG)~Direction, player_frame, sum)
direction_averages$xBA <- direction_averages$xBA / direction_averages$PA
direction_averages$xSLG <- direction_averages$xSLG / direction_averages$PA
direction_averages$xISO <- direction_averages$xSLG - direction_averages$xBA
direction_averages$BA <- direction_averages$BA / direction_averages$PA
direction_averages$SLG <- direction_averages$SLG / direction_averages$PA
direction_averages$ba_dif <- direction_averages$xBA - direction_averages$BA
direction_averages$slg_dif <- direction_averages$xSLG - direction_averages$SLG


direction_averages <- subset(direction_averages, direction_averages$Direction >= -45 &
                             direction_averages$Direction <= 45)

direction_line <- ggplot(data = direction_averages, aes(x = Direction, y = ba_dif)) +
  geom_line()
direction_line

```



```
direction_line2 <- ggplot(data = direction_averages, aes(x = Direction, y = slg_dif)) +  
  geom_line()  
direction_line2
```



There appear to still be some issues right down the foul lines, but overall the differences appear to be random. Obviously an ideal model would have no variation, but I feel that the variance that does exist is largely noise due to the relatively small size of this dataset and not due to an external force that requires correction.

```
player_averages <- aggregate(cbind(xBA,xSLG,PA)~Player, player_frame, sum)
player_averages$xBA <- player_averages$xBA / player_averages$PA
player_averages$xSLG <- player_averages$xSLG / player_averages$PA
player_averages$xISO <- player_averages$xSLG - player_averages$xBA

player_averages <- subset(player_averages, player_averages$PA >= 30)
player_averages <- player_averages[order(-player_averages$xISO),]

print(head(player_averages, 20))
```

##	Player	xBA	xSLG	PA	xISO
## 433	Cabell, Elijah	0.5190033	1.1835110	56	0.6645077
## 2037	McCann, Kyle	0.4879431	1.1087055	69	0.6207624
## 3228	Vaughn, Andrew	0.5058775	1.1039205	74	0.5980430
## 269	Bleday, JJ	0.4890859	1.0823803	105	0.5932945
## 158	Barr, Cole	0.5074929	1.1000640	56	0.5925711
## 3156	Torkelson, Spencer	0.5075968	1.0912515	124	0.5836547
## 2555	Radcliff, Baron	0.5224448	1.0840558	62	0.5616110
## 254	Bishop, Hunter	0.4887589	1.0338769	102	0.5451180
## 3291	Washer, Jake	0.5616378	1.0908222	31	0.5291844
## 1470	Horanski, Luke	0.4333701	0.9619095	59	0.5285394

## 231	Berryhill, Luke	0.4201209	0.9475818	91	0.5274610
## 2120	Mendoza, Drew	0.5149345	1.0408046	87	0.5258700
## 528	Chavers, Parker	0.4049588	0.9270379	83	0.5220790
## 2054	McDaniel, Joel	0.4512774	0.9659491	57	0.5146717
## 1655	Kavadas, Niko	0.4624297	0.9747977	84	0.5123680
## 2920	Smith, Armani	0.4637686	0.9738385	93	0.5100698
## 2718	Sabato, Aaron	0.4814773	0.9907795	113	0.5093022
## 1845	Lloyd, Matthew	0.4592847	0.9511124	68	0.4918277
## 2861	Shinn, Ryan	0.5013147	0.9923014	90	0.4909866
## 295	Boone, Trevor	0.4808353	0.9712040	62	0.4903687

There is slight movement from our initial table to now, meaning the correction was made in the model. An initial glance at the top performers according to this metric looks promising. Top picks like Vaughn, Bleday, Torkelson, Bishop, and Sabato all appearing towards the top indicates that this is capturing some level of power production.

The purpose of this model isn't to estimate overall production, but rather to anticipate how much power a particular batter may hit for WHEN they make contact. In the future, this metric can be referenced when compiling scouting reports so that coaches and pitchers can know which opposing hitters are most likely to hit for power if they are able to square up on a pitch. It is not an all-encompassing offensive metric, but helps tell a part of the story of a hitter.